

# Reactive Traversal of Recursive Data Types (?!)

Francisco Sant'Anna  
Departamento de Informática  
— PUC-Rio, Brazil  
fsantanna@inf.puc-rio.br

Hisham Muhammad  
Departamento de Informática  
— PUC-Rio, Brazil  
hisham@inf.puc-rio.br

Johnicholas Hines  
Affiliation  
email@domain.com

## ABSTRACT

We propose a structured mechanism to traverse recursive data types incrementally, in successive reactions to input events. `traverse` is an iterator-like anonymous block that can be invoked recursively and suspended at any point, retaining the full state and stack frames alive. `traverse` is designed for the synchronous language CÉU, inheriting all of its concurrency functionality and safety properties, such as parallel compositions with orthogonal abortion, static memory management, and bounded reaction time and memory usage. We discuss three applications in the domains of incremental computation and control-oriented DSLs that contain reactive and recursive behavior at the same time.

## Categories and Subject Descriptors

D.3.3 [Programming Languages]: Language Constructs and Features

## General Terms

Design, Languages

## Keywords

Behavior Trees, Domain Specific Languages, Incremental Computation, Logo, Recursive Data Types, Structured Programming, Reactive Programming

## 1. INTRODUCTION

... CÉU [2, 3]

...

```
1 input void RESET; // declares an external event
2 var int v = 0; // variable shared by the trails
3 par do
4     loop do // 1st trail
5         await 1s;
6         v = v + 1;
7         _printf("v = %d\n", v);
8     end
9 with
10    loop do // 2nd trail
11        await RESET;
12        v = 0;
13    end
14 end
```

Figure 1: Introductory example in CÉu.

## 2. CÉU

CÉU is a concurrent and reactive language in which the lines of execution, known as *trails*, react all together continuously and in synchronous steps to external stimuli. The introductory example in Figure 1 defines an input event `RESET` (line 1), a shared variable `v` (line 2), and starts two trails with the `par` construct (lines 3-14): the first (lines 4-8) increments variable `v` on every second and prints its value on screen; the second (lines 10-13) resets `v` on every external request to `RESET`. CÉU is tightly integrated with *C* and can access libraries of the underlying platform directly by prefixing symbols with an underscore (e.g., `_printf(<...>)`, in line 7).

In the synchronous model of CÉU, a program reacts to an occurring event completely before handling the next. A reaction represents a logical instant in which all trails awaiting the occurring event awake and execute, one after the other, until they await again or terminate. During a reaction, the environment is invariant and does not interrupt the running trails. If multiple trails react to the same event, the scheduler employs lexical order, i.e., the trail that appears first in the source code executes first. For this reason, programs are deterministic even in the presence of side effects in concurrent lines of execution. To avoid infinite execution for reactions, CÉU ensures that all loops contain `await` statements [2].

### 2.1 Recursive Data Types

The `data` construct in CÉU provides a safer alternative to C's `struct`, `union`, and `enum` definitions. A `data` entry declares either a non-recursive structure containing a set of mutable fields or a tagged union. A tagged union consists of a set of

```

1 data List with
2   tag NIL;
3 or
4   tag CONS with
5     var int head;
6     var List* tail;
7   end
8 end
9
10 do
11   pool List[1] lst1;
12   lst1 =new List.CONNS(10,
13     List.CONNS(20,
14       List.CONNS(30,
15         List.NIL()));
16   _printf("%d, %d\n", lst1:CONS.head,
17     lst1:CONS.tail:NIL);
18   // prints 10, 1
19 end
20
21 do
22   pool List[] lst2;
23   lst2 =new List.CONNS(10,
24     List.CONNS(20,
25       List.CONNS(30,
26         List.NIL()));
27   lst2:CONS.tail =new List.CONNS(50, List.NIL());
28   _printf("%d\n", lst2:CONS.tail:CONS.head);
29   // prints 50 (20 and 30 have been freed)
30 end

```

**Figure 2: A recursive List data type definition (lines 1–8) and uses (lines 10–18 and 20–28).**

tag declarations, each of which may be a bare tag or contain mutable fields. If any of the tag declarations refers to the data type being declared, we have a recursive data type. In this case, the first tag of the tagged union must be a bare tag, and it will act as the union’s null type: in CÉU, every recursive data type is an option type.

Figure 2 illustrates the recursive List data type, declared as a tagged union. The first tag, NIL (line 2), represents the empty list and is the union’s null type. The second tag, CONS, holds a value in its field head and a pointer to the rest of the list in the field tail (lines 4–7).

All memory allocated by CÉU constructs is managed by lexically-scoped memory pools. The pool keyword declares a memory pool of a given size and a reference to a root object. In line 11, we declare a pool of List objects of size 1, identified by root reference lst1, scoped by the do block in lines 10–19. The declaration also implicitly initializes the root to the null tag of the associated data type (i.e., List.NIL).

Then, in lines 12–15, we use the =new construct, which performs allocation and assignment: it attempts to dynamically allocate a list of three elements (using three List.CONNS constructors in the assignment *r-value*), inferring the destination memory pool based on the assignment’s *l-value* (i.e. lst1).

Since the pool has size 1, only the allocation of first element succeeds, with the failed allocations returning the null tag for this type (i.e., List.NIL). The print command (line 16) outputs “10, 1”: the head of the first element (the operator ‘.’ is equivalent to C’s ‘->’) and a true value for the NIL check of the second element.

```

1 pool List[3] l = new List.CONNS(1,
2   List.CONNS(2,
3     List.CONNS(3,
4       List.NIL())));
5
6 var int sum =
7   traverse e in l do
8     if e:NIL then
9       escape 0;
10    else
11      var int sum_tail = traverse e:CONS.tail;
12      escape sum_tail + e:CONS.head;
13    end
14  end;
15 _printf("sum = %d\n", sum);

```

**Figure 3: Calculating the sum of a list.**

In the second block (lines 21–30), we declare the lst2 pool with an unbounded memory limit (i.e., List[] in line 22). Now, the three-element allocation succeeds (lines 23–26). Then, we mutate the tail of the first element to point to a newly allocated element in the same pool, which also succeeds (line 27). The print command (line 28) outputs “50”, displaying the head of the new second element. In the moment of the mutation, the old subtree (containing values “20” and “30”) is completely removed from memory. Finally, the end of the block (line 30) deallocates the pool along with all of its elements.

In CÉU, recursive data types have a number of restrictions. Given that mutations deallocate whole subtrees, data types cannot represent general graphs: they must represent tree-like structures. Elements in different pools cannot be mixed; and pointers to subtrees (i.e., weak references) must be observed via the watching construct, as they can be invalidated at any time (to be discussed in Section 2.2).

## 2.2 Traversing Data Types

CÉU introduces a structured mechanism to traverse data types. The traverse construct integrates well with the synchronous execution model, supporting nested control compositions, such as await and all par variations. It also preserves explicit lexical scopes with static memory management.

We begin by showing the flavor of the construct through an example. The code in Figure 3 creates a list (lines 1–4) and traverses it to calculate the sum of elements (lines 6–15). The traverse block (line 7) starts with the element e pointing to the root of the list l. The escape statement (lines 9 and 12) returns a value to be assigned to the sum (line 6). A NIL list has sum=0 (lines 8–9). A CONS list needs to calculate the sum of its tail recursively, invoking traverse again (line 11), which will create a nested instance of the enclosing traverse block (lines 7–14), now with e pointing to the e:CONS.tail. When used without event control mechanisms, as in this simple example, a traverse block is equivalent to an anonymous closure called recursively.

The traverse construct does not simply amount to an anonymous recursive block, however. It is designed to take into account the event system and memory management discipline of the language. As such, it is an abstraction defined in terms of more fundamental CÉU features: *organisms*, which

```

1 pool List[] l = <...>; // 1, 2, 3
2 var int sum =
3   traverse e in l do
4     if e:NIL then
5       escape 0;
6     else
7       watching e do
8         _printf("me = %d\n", e:CONS.head);
9         await 1s;
10        var int sum_tail = traverse e:CONS.tail;
11        escape sum_tail + e:CONS.head;
12      end
13    end
14  end;
15 _printf("sum = %d\n", sum);
16

```

**Figure 4: Calculating the *sum* of a list, one element each second.**

are objects with their own parallel trail of execution, akin to Simula objects; and orthogonal abortion, which handles cancellation of trails maintaining memory consistency [3].

Three aspects make `traverse` fundamentally different from an anonymous recursive function. First, each `traverse` call spawns a new anonymous organism, launching a new parallel trail of execution (as opposed to stacking a new frame in the current trail). Second, traversal is declared in terms of a specific memory pool. Therefore, for bounded pools (e.g., `List[3]` 1), we can infer at compile time the maximum traversal depth. Third, the execution body of a `traverse` block is implicitly wrapped by a concurrency construct that watches for mutations of the current node. In practice, this means that it reacts consistently if another trail of execution modifies the data structure being traversed.

To illustrate these differences, Figure 4 extends the body of the previous example with reactive behavior. For each recursive iteration, `traverse` prints the current `head` (line 8) and awaits 1 second before traversing the `tail` (lines 9–10). In CÉU, all accesses to pointers that cross `await` statements must be protected with `watching` blocks [3]. This ensures that if side effects occurring in parallel affect the pointed object, no code uses stale pointers because the whole block is aborted. In the example (lines 7–12), if the list is mutated during that 1 second and the specific element is removed from memory, we simply ignore the whole subtree and return 0.

The `traverse` construct allows us to enforce bounded execution time, by performing a limited number of steps, each of them a separate synchronous reaction. This can be asserted by verifying that the structure of the recursive steps converge to the base cases, or simply by using a bounded memory pools, which allows us to limit the maximum number of steps to the size of the pool. Enforcing execution limits is an important requirement for constrained and real-time embedded systems, which is the original application domain of CÉU [2].

### 3. APPLICATIONS

In this section, we present three applications that explore the reactive and incremental nature of the `traverse` construct. We start with standard *Behavior Trees* used in video games

for AI modeling [?], and extend them with parallel compositions. Then, we show a *Logo Turtle* [?] that can execute commands in parallel (e.g., `move` and `rotate`) and also incorporates a dynamic queue of commands issued concurrently with the running program. Finally, we implement *Gray code* generation [?] to illustrate how `traverse` can also be used for more general recursive algorithms without an associated data type.

How many applications will we present?

- behavior trees (DSL)
- logo turtle (DSL)
- turtle stream (incremental)
- gray code gen (incremental)

I suggest gray code as the last one because it does not use a data type. We could show as an example of flexibility of `traverse`. BTs before turtle because they are simpler and because the stream example is the most complex and should follow the logo turtle.

#### 3.1 Behavior Trees

The term *Behavior Trees* denotes a family of DSLs used for Game AI [?]. The term is loose, because different games use different languages, but generally it indicates an interpreted domain-specific language for creature behavior that includes at least sequence and selection combinators, and which are “ticked” periodically [?].

The semantics of the sequence combinator can be understood as short-circuit evaluation of a conjunction; the `SEQ` node ticks its left subtree until it finishes, and if it finishes successfully, ticks its right subtree until it finishes. The semantics of the selection combinator can be understood as short-circuit evaluation of an alternation; the `SEL` node ticks its left subtree until it finishes, and if it did not finish successfully, ticks its right subtree until it finishes.

This skeleton, augmented with leaves that test properties, set properties, perform animations and sounds, and other custom combinators, can be preferable to finite state machines (hierarchical, augmented, or otherwise) for authoring Game AI.

Ceu’s parallel features make implementing a parallel combinator for behavior trees much easier.

Most “standard” BTs I found have these “sequence” and “selector” composite nodes. I thought about starting with this one and maybe expanding it with a “parallel/or” further.

Two things to discuss: (1) how the implementation of the interpreter is straightforward; (2) that leaf nodes are not restricted to a “tick” callback and can actually execute arbitrary code in CÉU.

```

1 data BTree with
2   tag NIL;
3 or
4   tag SEQ with
5     var BTree* first;
6     var BTree* second;
7   end
8 or
9   tag SEL with
10    var BTree* first;
11    var BTree* second;
12  end
13 or
14   tag LEAF with
15     var Leaf leaf;
16   end
17 end

```

Figure 5: A standard behavior tree with *sequence* and *selector* composite nodes.

```

1 class BTreeInterpreter with
2   pool BTree[] & btree;
3 do
4   var int ret =
5     traverse t in btree do
6       if t:SEQ then
7         var int ret1 = traverse t:SEQ.first;
8         if ret1 > 0 then
9           var int ret2 = traverse t:SEQ.second;
10          if ret2 > 0 then
11            escape ret1+ret2;
12          end
13        end
14        escape 0;
15      else if t:SEL then
16        var int ret = traverse t:SEQ.first;
17        if ret == 0 then
18          ret = traverse t:SEL.second;
19        end
20        escape ret;
21      else if t:LEAF then
22        var int ret =
23          do LeafHandler(t:LEAF.leaf);
24        escape ret;
25      end
26    end;
27  escape ret;
28 end
29
30 pool BTree[] btree = new
31   BTree.SEQ(
32     BTree.SEL(
33       BTree.LEAF(Leaf(0)),
34       BTree.LEAF(Leaf(1)),
35       BTree.LEAF(Leaf(2)));
36
37 var int ret = do BTreeTraverse(btree);
38
39 _printf("ret = %d\n", ret); // prints "3"

```

Figure 6: A straightforward interpreter for the standard behavior tree of Figure 5 and a sample tree to execute.

```

1 class LeafHandler with
2   var Leaf& leaf;
3 do
4   // TODO: what to show here?
5   escape leaf.v;
6 end

```

Figure 7: A leaf node with complex behavior.

## 3.2 Logo Turtle

Our second example is an interpreter for a simple variant of the classic Logo turtle-graphics interpreter [1], which extends the Logo paradigm with a CÉU-like parallel execution construct. In our variant, we can instruct the turtle to move and rotate in parallel, tracing curves. We declare a data type which defines our abstract syntax, with each tag representing one of the supported Logo commands. A tree of nodes represents a program, and the interpreter is implemented as a traversal of this tree. The aim of this example is to demonstrate parallel traversal.

Figure 8 presents the data type `Command`, which specifies the abstract syntax of our Logo variant. As in traditional Logo, commands can be listed in sequence to be executed one after the other (represented through a chain of `SEQ` nodes), and commands can be repeated a number of times (denoted through a `REPEAT` node). Our variant includes `MOVE` and `ROTATE` nodes to move the turtle, but these are specified differently from traditional Logo: here, they take as arguments the speed at which they should affect the turtle. For example, a `Command.MOVE(50)` node directs the turtle to move at the speed of 50 pixels per second, indefinitely. The only way to make the turtle stop moving or rotating is through two CÉU-like extensions added to our Logo variant: `AWAIT` and `PAROR`. `AWAIT` simply awaits a given number of milliseconds. `PAROR`, modeled after the CÉU construct `par/or`, launches two commands in parallel, and aborts both of them as soon as one of them finishes. For example, the following construct would make the turtle move along a semicircle:

```

Command.PAROR(
  Command.AWAIT(1000),
  Command.PAROR(
    Command.MOVE(50),
    Command.ROTATE(180)))

```

Figure 9 depicts the interpreter. It is implemented as the `Interpreter` organism (declared with the `class` keyword). It holds as attributes a reference to the AST of commands (`cmds`, line 2) and a reference to a `Turtle` object which implements the UI. The execution body of the organism contains the `traverse` construct which runs the interpreter (lines 5–35).

We have then a test for each kind of tag in the data type. In lines 7–9, `SEQ` is handled by traversing each of its child commands, in sequence: the second invocation of `traverse` in that block (line 9) only runs after the first one (line 8) finishes. In lines 11–14, `REPEAT` is handled by traversing its command the specified number of times.

`MOVE` is handled in lines 16–18 by spawning a new organism called `TurtleMove`, which launches a separate trail of execution. The implementation of `TurtleMove` (not shown) updates the coordinates of the `turtle` instance it got as a parameter in its constructor. The implementation of `ROTATE` (lines 20–22) is similar.

In lines 24–25, `AWAIT` is implemented by simply causing the current trail of execution of the interpreter to await the given amount of time. Finally, `PAROR` (lines 27–32) uses the `par/or`

```

1 data Command with
2   tag NOTHING;
3 or
4   tag SEQ with
5     var Command* one;
6     var Command* two;
7   end
8 or
9   tag REPEAT with
10    var int times;
11    var Command* command;
12  end
13 or
14   tag MOVE with
15     var int pixels;
16   end
17 or
18   tag ROTATE with
19     var int angle;
20   end
21 or
22   tag AWAIT with
23     var int ms;
24   end
25 or
26   tag PAROR with
27     var Command* one;
28     var Command* two;
29   end
30 end

```

Figure 8: DSL for a Logo turtle.

construct to traverse both subcommands at the same time. As per the semantics of `par/or`, as soon as one of the subtrees terminate its execution, the other one will be aborted.

Note that the entire interpreter block is surrounded by a `watching` construct (line 6). The C  U compiler enforces the presence of a guard, due to the use of the `cmd` pointer in code that spans multiple reactions. This ensures clean abortion in case the AST being interpreted is mutated by code running in another trail.

### 3.2.1 Enqueuing Commands

## 3.3 Gray Code Generation

## 4. RELATED WORK

...

What to compare against???

Found this: <http://arxiv.org/pdf/1104.2293.pdf>

Follow the references.

## 5. CONCLUSION

...

Limitations:

- high-order programming

- planar-only recursive data types

## 6. REFERENCES

- [1] S. Papert. *Mindstorms: Children, Computers, and Powerful Ideas*. Basic Books, Inc., New York, NY, USA, 1980.

```

1 class Interpreter with
2   pool Command[] & cmds;
3   var Turtle& turtle;
4 do
5   traverse cmd in cmds do
6     watching cmd do
7       if cmd:SEQ then
8         traverse cmd:SEQ.one;
9         traverse cmd:SEQ.two;
10
11       else/if cmd:REPEAT then
12         loop i in cmd:REPEAT.times do
13           traverse cmd:REPEAT.command;
14         end
15
16       else/if cmd:MOVE then
17         do TurtleMove(turtle,
18           cmd:MOVE.pixels);
19
20       else/if cmd:ROTATE then
21         do TurtleRotate(turtle,
22           cmd:ROTATE.angle);
23
24       else/if cmd:AWAIT then
25         await (cmd:AWAIT.ms) ms;
26
27       else/if cmd:PAROR then
28         par/or do
29           traverse cmd:PAROR.one;
30           with
31             traverse cmd:PAROR.two;
32           end
33         end
34       end
35     end
36 end

```

Figure 9: The turtle interpreter.

```

1 var int tot = 4;
2 var int[] vec = [0, 0, 0, 0];
3
4 par/or do
5   every VISIT do
6     _printf(" ");
7     loop i in tot do
8       _printf("%d ", vec[i]);
9     end
10    _printf("\n");
11  end
12 with
13   /*
14    * Syntax invented for numeric/non-data
15    * traversals (not implemented):
16    * - starts from idx=0
17    * - memory for at most [N] elements,
18    *   or [] unbounded.
19    * - recursive calls receive the next idx
20    */
21   traverse idx in [] do
22     if idx == tot then
23       await NEXT;
24     else
25       traverse idx + 1;
26       vec[idx] = 1 - vec[idx];
27       traverse idx + 1;
28     end
29   end
30 end

```

Figure 10: Gray code generation for tuples of size 4.

- [2] F. Sant'Anna et al. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM, 2013.
- [3] F. Sant'Anna et al. Structured Synchronous Reactive Programming with Céu. In *Proceedings of Modularity'15*, 2015. to appear.