

# Reactive Traversal of Recursive Data Types

Francisco Sant’Anna  
Departamento de Informática  
— PUC-Rio, Brazil  
fsantanna@inf.puc-rio.br

Hisham Muhammad  
Departamento de Informática  
— PUC-Rio, Brazil  
hisham@inf.puc-rio.br

Johnicholas Hines  
IDEXX Laboratories  
johnicholas.hines@gmail.com

## ABSTRACT

We propose a structured mechanism to traverse recursive data types incrementally, in successive reactions to external input events. `traverse` is an iterator-like anonymous block that can be invoked recursively and suspended at any point, retaining the full state and stack frames alive. `traverse` is designed for the synchronous language CÉU, inheriting all of its concurrency functionality and safety properties, such as parallel compositions with orthogonal abortion, static memory management, and bounded reaction time and memory usage. We discuss three applications in the domain control-oriented DSLs that contain reactive and recursive behavior at the same time.

## Categories and Subject Descriptors

D.3.3 [Programming Languages]: Language Constructs and Features

## General Terms

Design, Languages

## Keywords

Behavior Trees, Domain Specific Languages, Incremental Computation, Logo, Recursive Data Types, Structured Programming, Reactive Programming

## 1. INTRODUCTION

The facilities a given language offers for constructing data types have a direct impact on the nature of algorithms that programmers will write on that language. As an example, the aim for referential transparency in functional languages enforces data structures to be immutable. Under these constraints, one must avoid excessive memory copying through specialized algorithms [12].

In this paper, we discuss the design of recursive data types and an associated control facility for a language developed under a different set of constraints. CÉU [14, 15]

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.  
Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

```
1  input void RESET; // declares an external event
2  var int v = 0;    // variable shared by the trails
3  par do
4      loop do      // 1st trail
5          await 1s;
6          v = v + 1;
7          _printf("v = %d\n", v);
8      end
9  with
10     loop do      // 2nd trail
11         await RESET;
12         v = 0;
13     end
14 end
```

Figure 1: Introductory example in CÉu.

is an imperative, concurrent and reactive language in which lines of execution, known as *trails*, react together continuously and in synchronous steps to external stimuli. CÉU supports mutable data, with static memory and safe pointer manipulation. These features are incompatible with garbage-collected immutable data structures, as well as with general records with arbitrary pointers such as *structs* in C.

The solution to this problem is twofold, with data and control aspects. For data management, we introduce a restricted form of recursive data types that supports trees (but not general graphs). To control reactive behavior, we propose a structured mechanism that can traverse data types safely and incrementally, in successive reactions to events. After we present the design of these constructs, we discuss three applications in the domain of control-oriented DSLs.

## 2. CÉU CONSTRUCTS

The introductory example<sup>1</sup> in Figure 1 gives a general flavor of CÉU. It first defines an input event `RESET` (line 1), a shared variable `v` (line 2), and starts two trails with the `par` construct (lines 3-14): the first (lines 4-8) increments variable `v` on every second and prints its value on screen; the second (lines 10-13) resets `v` on every external request to `RESET`.

In the synchronous model of CÉU, a program reacts to an occurring event completely before handling the next. A reaction represents a logical instant in which all trails awaiting the occurring event awake and execute *atomically*, one after the other, until they await again or terminate. As a consequence, all consecutive operations to shared variable `v` in Figure 1 are atomic because reactions to events `1s` and

<sup>1</sup>A screencast of all examples in the paper is available at <https://vimeo.com/135297440>.

RESET can never interrupt each other. If multiple trails react to the same event, the scheduler employs lexical order, i.e., the trail that appears first in the source code executes first. For this reason, programs are deterministic even in the presence of side effects in concurrent lines of execution. To avoid infinite execution for reactions, CÉU ensures that all loops contain `await` statements [14].

## 2.1 Recursive Data Types

The `data` construct in CÉU provides a safer alternative to C’s `struct`, `union`, and `enum` definitions. Figure 2 illustrates the recursive `List` data type, declared as a tagged union (lines 1–5). The first tag, `NIL` (line 2), represents the empty list and is the union’s null type. The second tag, `CONS` (line 4), receives two values in the constructor: the field `int` `head` and the rest of the list in the recursive field `List` `tail`.

In the first block of the example (lines 7–16), we declare a `pool` of `List` objects of size 1 (line 8). All recursive data instances must reside in an explicit memory pool, which have static memory management based on its lexical scope (lines 7–16). A pool also represents the reference to the root instance, which is implicitly initialized to the `null` tag of the associated data type, i.e., `lst1` receives `List.NIL` (line 8). Then, we use the `=new` construct (lines 9–12) which performs allocation and assignment at the same time: it attempts to dynamically allocate a list of three elements (10, 20, and 30), inferring the destination memory pool based on the assignment’s *l-value* prefix (i.e. `lst1`). Since the pool has size 1, only the allocation of first element succeeds, with the failed subtree allocation returning the `null` tag (i.e., `List.NIL`). The print command (lines 13–14) outputs “10, 1”: the `head` of the first element and a true value for the `NIL` check of the second element. Finally, the end of the block (line 16) deallocates the pool along with all elements inside it.

In the second block (lines 18–24), we declare the `lst2` pool with an unbounded memory limit (i.e., `List[]` in line 19). Now, the three-element allocation succeeds (line 20)<sup>2</sup>. Then, we mutate the tail of the first element to point to a newly allocated element in the same pool, which also succeeds (line 21). In the moment of the mutation, the old subtree (containing values “20” and “30”) is completely removed from memory. The print command (line 22) outputs “50”, displaying the `head` of the new second element. Again, the end of the block (line 24) deallocates the pool along with all of its elements.

In CÉU, recursive data types have a number of restrictions. Given that mutations deallocate whole subtrees, data types cannot represent general graphs with cycles, but only tree-like structures. Also, elements in different pools cannot be mixed. Finally, pointers to subtrees (i.e., weak references) must be observed via the `watching` construct, as they can be invalidated at any time (to be discussed in Section 2.2).

## 2.2 Traversing Data Types

CÉU introduces a structured mechanism to traverse recursive data types. The `traverse` construct integrates with the synchronous execution model of CÉU and supports nested control compositions, such as `await` and all `par` variations.

The code in Figure 3 creates a list (line 1) and traverses it to calculate the sum of elements (lines 3–11). The `traverse`

```

1  data List with
2    tag NIL ();
3  or
4    tag CONS (int head, List tail);
5  end
6
7  do
8    pool List[1] lst1;
9    lst1 = new List.CONS(10,
10                       List.CONS(20,
11                                List.CONS(30,
12                                           List.NIL()));
13    _printf("%d, %d\n", lst1.CONS.head,
14              lst1.CONS.tail.NIL);
15    // prints 10, 1
16  end
17
18  do
19    pool List[] lst2;
20    lst2 = new CONS(10, CONS(20, CONS(30, NIL())));
21    lst2.CONS.tail = new CONS(50, NIL());
22    _printf("%d\n", lst2.CONS.tail.CONS.head);
23    // prints 50 (20 and 30 have been freed)
24  end

```

Figure 2: A recursive `List` data type definition (lines 1–5) with uses (lines 7–16 and 18–27).

```

1  pool List[3] lst = <...>; // [10, 20, 30]
2
3  var int sum =
4    traverse e in lst do
5      if e:NIL then
6        escape 0;
7      else
8        var int sum_tail = traverse e:CONS.tail;
9        escape sum_tail + e:CONS.head;
10     end
11  end;
12
13 _printf("sum = %d\n", sum);
14 // prints 60

```

Figure 3: Calculating the *sum* of a list.

block (lines 4–11) starts with the element `e` pointing to the root of the list `lst`. The `escape` statement (lines 6 and 9) returns a value to the enclosing assignment to `sum` (line 3). A `NIL` list<sup>3</sup> has `sum=0` (lines 5–6). A `CONS` list needs to calculate the `sum` of its tail recursively, invoking `traverse` again, which will create a nested instance of the enclosing `traverse` block (lines 4–11), now with `e` pointing to `e:CONS.tail` (line 8). Only after its complete subtree is traversed recursively that it adds its `head` and returns (line 9).

When used without event control mechanisms, as in this simple example, a `traverse` block is equivalent to an anonymous closure called recursively. However, `traverse` complies with the event system and memory management discipline of CÉU and is an abstraction defined in terms of a more fundamental concept [15]: *organisms*, which are objects with concurrent trails of execution (akin to Simula [3]); and orthogonal abortion, which handles cancellation of trails maintaining memory consistency. Figure 4 depicts the expansion of the `traverse` construct.

The example in *CODE-1* of Figure 4 extends the body of the previous example in Figure 3 with reactive behavior. Now, for each recursive iteration, we print the current `head` (line 9) and `await` 1 second (line 10) before traversing the `tail` (line 11). Note that the last iteration of `traverse` will

<sup>2</sup>To save space, in the next examples we omit the data type prefix in tags (e.g., `List.CONS` becomes `CONS`).

<sup>3</sup>The operator ‘:’, as in `e:NIL`, is equivalent to C’s ‘->’.

```

1  pool List[3] lst = <...>; // [10, 20, 30]
2
3  var int sum =
4    traverse e in lst do
5      if e:NIL then
6        escape 0;
7      else
8        watching e do
9          _printf("me = %d\n", e:CONS.head);
10         await 1s;
11         var int sum_tail = traverse e:CONS.tail;
12         escape sum_tail + e:CONS.head;
13       end
14     end
15   end
16 end;
17
18 _printf("sum = %d\n", sum);
19 // prints 60
20
21
22
23
24
25
26
27
28
29
30
31
32
33 .

```

CODE-1: Original code (with traverse)

```

1  pool List[3] lst = <...>; // [10, 20, 30]
2
3  class Frame with
4    pool Frame[3]& frames;
5    var Frame& parent;
6    pool List[3]* e;
7  do
8    watching this.parent do
9      if e:NIL then
10        escape 0;
11      else
12        watching e do
13          _printf("me = %d\n", e:CONS.head);
14          await 1s;
15          var Frame* frame = spawn Frame(this.frames,
16                                         this,
17                                         e:CONS.tail);
18                                in this.frames;
19          var int sum_tail = await *frame;
20          escape sum_tail + e:CONS.head;
21        end
22      end
23    end
24  end
25  escape 0;
26 end
27
28 pool Frame[3] frames;
29 var Frame* frame = spawn Frame(frames, this, lst)
30                   in frames;
31 var int sum = await *frame;
32
33 _printf("sum = %d\n", sum); // prints 60

```

CODE-2: Expanded code (without traverse)

Figure 4: Calculating the *sum* of a list, one element each second. The *traverse* construct is a syntactic sugar that can be “desugared” with explicit organisms.

wait for one second before printing “30”, with all previous iterations blocked and retaining their full state of execution. Furthermore, this code can be part of a larger program with other trails in parallel, all of which remain reactive during the incremental traversal.

CÉU enforces at compile time that all accesses to a pointer that cross *await* statements are protected with an enclosing *watching* block, which aborts its nested block if the referred object is released from memory [15]. This ensures that if concurrent side effects affect the pointed object, no code uses the stale pointer, because the whole block is aborted. With the protection of the *watching* block (lines 8–13), if the element *e* (line 8) is released from memory due to a mutation in the list during the awaiting period (line 10), we simply ignore the whole subtree and return 0 (line 14).

CODE-2 is the equivalent expansion of CODE-1 without the *traverse* construct. Because it contains concurrency constructs (i.e., *await* and *watching*), the body of the *traverse* (CODE-1: 5–15) must be abstracted in an organism of the *Frame* class (CODE-2: 3–26), which is analogous to a “stack frame” of standard programming languages with sub-routines. Likewise, the pool of frames (CODE-2: 28) is analogous to a runtime “call stack”. We limit the number of stack frames to match the maximum number of elements to traverse (CODE-1: 1 and CODE-2: 1,28). Therefore, to “call” the first *traverse* iteration, we dynamically *spawn* a *Frame* instance into the *frames* pool (CODE-2: 29–30). Then, we immediately *await* the termination of this *frame* (CODE-2: 31). Only after the *traverse* rolls back and clears the whole call stack that we acquire the *sum* and print it (CODE-2:

31–34).

A *Frame* receives three arguments in the constructor (CODE-2: 4–6): a reference to a pool (to recursively *spawn* new frames); a reference to its parent frame (to handle abortion); and a pointer to the subtree of the data type (to be able to manipulate it). The *Frame* constructor for the first call (CODE-2: 29–30) receives the static pool of frames, the running organism as the parent (i.e., *this*), and the original tree to traverse (CODE-1: 4). The *Frame* constructor for recursive calls (CODE-2: 15–18) receives the same pool of frames, the current stack frame as parent, and the original subtree in the recursive invocation (CODE-1: 11).

The *Frame* body (CODE-1: 8–25) always aborts with the parent frame termination. Given that a *traverse* body can possibly execute (and terminate) trails running concurrently with the recursive invocation, the enclosing *watching* guarantees that the hierarchy in the call stack is preserved (i.e., that there are no orphan frames executing). The remaining code is almost the same in the original *traverse* body and in the *Frame* body (CODE-1: 5–15 and CODE-2: 9–23), with the exception of the recursive invocation explained above (CODE-2: 15–18).

As the expansion illustrates, three aspects make *traverse* fundamentally different from recursive function calls:

1. Each *traverse* invocation spawns a new organism which can execute concurrently with other parts of the application. Also, each organism itself can invoke multiple concurrent trails, adding more complexity in comparison to standard functions [15].
2. A *traverse* is attached to a specific lexically-scoped

memory pool for a data structure. Therefore, we can infer at compile time the maximum traversal depth if the data is bounded (e.g., `List[3] lst`). Enforcing execution limits is an important requirement for constrained and real-time embedded systems, which is the original application domain of CÉU [14]. In addition, given that the pool of frames is expanded in the same lexical scope, when the associated data goes out of scope, all stack frames are automatically aborted [15].

3. The execution body of a `traverse` block is implicitly wrapped by a concurrency construct that watches for mutations of the current node. In practice, this means that it reacts consistently if another trail of execution modifies the data structure being traversed.

We believe that the `traverse` construct, more than a simple convenience, considerably reduces the complexity of programs, handling a complex hierarchy of behaviors associated with recursive data types automatically.

### 3. APPLICATIONS

In this section, we present three applications that explore the reactive and incremental nature of the `traverse` construct. We start with *Behavior Trees* used in video games for AI modeling [9, 6]. Then, we show a *Logo Turtle* [13] that can execute commands in parallel (e.g., `move` and `rotate`). Finally, we extend the Turtle example with a dynamic and concurrent queue of commands that can affect the running program.

#### 3.1 Behavior Trees

*Behavior Trees* are a family of DSLs used for game AI. The term is loose, because different games use different languages. For our purposes it indicates an interpreted domain-specific language for concurrent creature behavior that includes sequence (`SEQ`) and selection (`SEL`) combinators.

The `SEQ` can be understood as short-circuit evaluation of an ‘and’, while the `SEL` corresponds to an ‘or’. This skeleton is extensible with leaves to test properties, set properties, perform animations and sounds, etc., and is an effective alternative to finite state machines for authoring game AI. However, because the evaluation of each tree extends across many frames, writing behavior tree nodes and leaves in other languages is an exercise in stack-ripped event-driven programming [10]. By lowering the barrier to writing custom nodes and leaves, CÉU lightweight event control mechanisms make behavior trees more usable.

Figure 5 describes a generic grammar for behavior trees (lines 1–9). The `SEQ` and `SEL` tags (lines 4 and 6) are recursive and behave as described above. The `LEAF` tag (line 8) receives a reference to a `Leaf` data type, which is defined externally, carrying opaque values that are specific to the application domain. The interpreter for trees is abstracted in a class definition (lines 11–37) and receives the tree to traverse as the argument (line 12). The body acquires the return status of the traversal (line 14) and returns it as its final result (lines 36). The traversal takes into account the semantics of the composite nodes as described before. For the `SEQ` tag (lines 16–22), we traverse the `first` subtree (line 17) and only if it succeeds, we traverse the `second` subtree (line 21). For the `SEL` tag (lines 23–29), we traverse the `first` subtree (line 24) and only if it fails, we traverse the `second` subtree (line 28). Finally, the `LEAF` tag does the real work and is domain specific. The `do` `Class` syntax (line 31–32) creates

```

1  data BTree with
2    tag NIL ();
3  or
4    tag SEQ (BTree first, BTree second);
5  or
6    tag SEL (BTree first, BTree second);
7  or
8    tag LEAF (Leaf& leaf);
9  end
10
11 class BTreeInterpreter with
12   pool BTree[]& btree;
13 do
14   var int ret =
15     traverse t in btree do
16       if t:SEQ then
17         var int ok = traverse t:SEQ.first;
18         if ok == 0 then
19           escape ok;
20         end
21         ok = traverse t:SEQ.second;
22         escape ok;
23       else/if t:SEL then
24         var int ok = traverse t:SEL.first;
25         if ok != 0 then
26           escape ok;
27         end
28         ok = traverse t:SEL.second;
29         escape ok;
30       else/if t:LEAF then
31         var int ret =
32           do LeafHandler(t:LEAF.leaf);
33         escape ret;
34       end
35     end;
36   escape ret;
37 end

```

Figure 5: A simple generic grammar of behavior trees with *sequence* and *selector* composite nodes with a straightforward interpreter.

an anonymous and lexically scoped organism of the referred class and awaits its termination. The organism itself can contain any valid code in CÉU and execute for an arbitrary amount of time [15]. The tree traversal hangs and waits for its return status (line 33), which is the expected behavior given that both `SEQ` and `SEL` nodes execute their children in sequence.

As an example of a domain, the *blocks world* is a classical planning domain in AI [17]. The tree in Figure 6 is based on the output from a Contingent-FF benchmark that extends the blocks domain with sensor actions [7]. Here, we want to achieve an ABC stack and assume we pinned down the situation to two possibilities illustrated in the figure. We use a `SEL` node (line 3) with a sensor leaf (line 4) to decide which strategy is appropriate: If `c` is not sensed on top of the table, we first move it to the table (line 5). Then, we stack `B` on top of `A`, and `C` on top of `B` (lines 6–7).

The example illustrates how the behavior tree can exhibit goal-directed behavior specified directly by domain designers. The `traverse` feature allows behavior authored at runtime, which is one of the crucial features of behavior trees. Alternatively, `traverse` allows CÉU to consume at runtime the output of CPU-intensive algorithms such as planners.

#### 3.2 Logo Turtle

Our second example is an interpreter for a simple variant of the classic Logo turtle-graphics interpreter [13]. The aim of this example is to demonstrate parallel traversal. In our

```

1 pool BTree[] btree =
2   new SEQ(SEQ,
3     SEL(
4       LEAF (SENSE_ON_TABLE(C)),
5       LEAF (MOVE_BLOCK_TO_TABLE(C)),
6       LEAF (MOVE_BLOCK_TO_BLOCK(B, A)),
7       LEAF (MOVE_BLOCK_TO_BLOCK(C, B)));

```




Figure 6: A blocks world behavior tree.

variant, we can instruct the turtle to move and rotate in parallel, tracing curves.

Figure 7 presents the data type `Command` (lines 1–15), which specifies the abstract syntax of our Logo variant. As in traditional Logo, commands can execute in sequence through the `SEQ` tag (line 4), and can also repeat a number of times through the `REPEAT` tag (line 6). Our variant extends the `MOVE` and `ROTATE` commands to take as arguments the speed at which they should affect the turtle (lines 8,12). For example, a `Command.MOVE(300)` node directs the turtle to move at the speed of 300 pixels per second, indefinitely. Therefore, the only way to make the turtle stop moving or rotating is through two CÉU-like extensions added to our Logo variant: The `AWAIT` tag (line 12) simply awaits a given number of milliseconds. The `PAROR` tag (line 14), modeled after the CÉU construct `par/or`, launches two commands in parallel, and aborts both of them as soon as one of them finishes. As an example, the program (lines 54–60) makes the turtle to move along a semicircle.

The interpreter for the commands is also abstracted in a class definition (lines 17–52). It holds as attributes a reference to a `Turtle` object (which implements the UI) and a reference to the commands (lines 18–19). The execution body of the class uses the `traverse` construct to interpret the commands (lines 21–51). The `SEQ` tag (lines 23–25) traverses each of its child commands in sequence (in contrast with the `BTreeInterpreter`, it does not handle failures). The `REPEAT` tag (lines 27–30) traverses its command the specified number of times. The `MOVE` and `ROTATE` tags (lines 32–38) rely on predefined classes of organisms to update the position and orientation of the `turtle` received as argument in the constructor (line 18). The `AWAIT` tag (lines 40–41) simply causes the current trail of execution of the interpreter to await the given amount of time. Finally, the `PAROR` tag (lines 43–48) uses the `par/or` construct to traverse both subcommands at the same time. As per the semantics of `par/or`, as soon as one of the subtrees terminate its execution, the other one is aborted.

Note that the entire interpreter block is surrounded by a `watching` construct (line 22). As discussed in Section 2.2, the CÉU compiler enforces the presence of a guard, due to the use of the `cmd` pointer in code that spans multiple reactions. This ensures clean abortion in case the AST is mutated by code running in other trails.

### 3.3 Enqueuing Commands

All examples so far create a fixed tree that does not vary during traversal. Figure 8 extends the Turtle application with a queue of pending commands to execute after the running commands terminate.

```

1 data Command with
2   tag NOTHING ();
3 or
4   tag SEQ (Command first, Command second);
5 or
6   tag REPEAT (int times, Command command);
7 or
8   tag MOVE (int pixels);
9 or
10  tag ROTATE (int angle);
11 or
12  tag AWAIT (int ms);
13 or
14  tag PAROR (Command first, Command second);
15 end
16
17 class CommandInterpreter with
18   var Turtle& turtle;
19   pool Command[]* cmds;
20 do
21   traverse cmd in cmds do
22     watching cmd do
23       if cmd:SEQ then
24         traverse cmd:SEQ.first;
25         traverse cmd:SEQ.second;
26
27       else/if cmd:REPEAT then
28         loop i in cmd:REPEAT.times do
29           traverse cmd:REPEAT.command;
30         end
31
32       else/if cmd:MOVE then
33         do TurtleMove(turtle,
34           cmd:MOVE.pixels);
35
36       else/if cmd:ROTATE then
37         do TurtleRotate(turtle,
38           cmd:ROTATE.angle);
39
40       else/if cmd:AWAIT then
41         await (cmd:AWAIT.ms) ms;
42
43       else/if cmd:PAROR then
44         par/or do
45           traverse cmd:PAROR.first;
46         with
47           traverse cmd:PAROR.second;
48         end
49       end
50     end
51   end
52 end
53
54 pool Command[] cmds =
55   new PAROR (
56     AWAIT(1000),
57     PAROR(MOVE(300), ROTATE(180)));
58
59 var Turtle turtle;
60 do CommandInterpreter(turtle, cmds);

```

Figure 7: Grammar, interpreter, and sample program for a Logo turtle DSL.

We define a new `Queue` data type in *CODE-3*: The `ROOT` tag (lines 4–6) has a `running` subtree of commands, a `waiting` queue of pending commands to execute afterwards, and a `tmp` node to allow in-place manipulation of the tree (to be discussed further). The `ITEM` tag (lines 8–9) represents a queue item and contains a `cmds` subtree with the actual commands to execute, and a `prv` queue item pointing to an older item (i.e., the queue is in reverse order). As Figure 9 illustrates in box 0, a queue instance has a single `ROOT` node with linked lists of `ITEM` nodes in the `running` and `waiting` fields. Except when creating a new command, the `tmp` field is always `NIL`.

We define the queue traversal in *CODE-4*. The `ROOT`

```

data Queue with
  tag NIL ();
or
  tag ROOT (Queue running,
             Queue waiting,
             Queue tmp);
or
  tag ITEM (Command cmds,
            Queue prv);
end

```

CODE-3: Queue type

```

1  traverse qu in queue do
2    watching qu do
3      if qu:ROOT then
4        loop do
5          par/and do
6            traverse qu:ROOT.running;
7            with
8              await qu:ROOT.waiting;
9            end
10           qu:ROOT.running =
11             qu:ROOT.waiting;
12           qu:ROOT.waiting =
13             new ITEM(NOTHING(), NIL());
14         end
15       else/if qu:ITEM then
16         traverse qu:ITEM.prv;
17         do CommandInterpreter(
18           turtle, qu:ITEM.cmds);
19       end
20     end
21   end

```

CODE-4: Queue traversal

```

1  input (char*,int,int) ENQUEUE;
2  every (cmd,vel,time) in ENQUEUE do
3    if _strcmp(cmd,"move")==0 then
4      queue.ROOT.tmp =
5        new ITEM(
6          NOTHING(),
7          ITEM(
8            PAROR(
9              MOVE(vel),
10             AWAIT(time)),
11            NIL()));
12    else/if _strcmp(cmd,"rotate")==0 then
13      <...> // analogous to the MOVE above
14    end
15    queue.ROOT.tmp.ITEM.prv.ITEM.prv =
16      queue.ROOT.waiting.ITEM.prv;
17    queue.ROOT.waiting = queue.ROOT.tmp;
18  end
19
20
21

```

CODE-5: Enqueuing commands

Figure 8: Queue extension for the Turtle DSL of Figure 7.

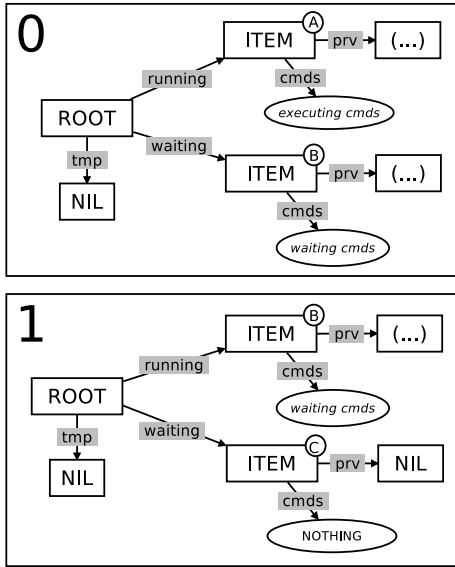


Figure 9: Swapping waiting and running commands.

traversal (lines 3–14) is a continuous loop that executes the **running** subtree and swaps it with the **waiting** queue on termination. The **par/and** (lines 5–9) ensures that that the swap only occurs after the current **running** commands terminate (line 6) *and* something (in parallel) mutates the **waiting** subtree (line 8), meaning that the queue is no longer empty. The swapping process (lines 10–13) is illustrated in Figure 9:

- The initial state (box 0) assumes pre-existing **running** and **waiting** items.
- Lines 10–11 assign the **waiting** subtree (marked ⑥) to the **running** field, releasing the old subtree (④)). The **waiting** field is automatically set to NIL.
- Lines 12–13 assign a new neutral ITEM (box 1: ③) to the **waiting** field, finishing the swapping operation.

After the swapping process, the loop restarts and traverses the new **running** commands (lines 4–14). The ITEM traversal (lines 15–18) is straightforward: first we traverse the previous item (line 16), and then we reuse the **CommandInterpreter**

class of Figure 7 to traverse the commands (line 17–18).

Even though this example mutates the **running** field only *after* its traversal terminates (lines 10–11), it is safe to do an arbitrary mutation at any time. Note that the compiler enforces the use of the **watching** construct (line 2) which encloses the running turtle interpreter (lines 17–18). Hence, if its enclosing ITEM (line 15) is mutated, the **watching** will awake and abort the interpreter which runs inside its lexical scope.

The enqueuing of new commands is depicted in *CODE-5*. The external input event **ENQUEUE** (line 1) accepts *move* and *rotate* commands with an associated velocity and time (i.e., “char\*,int,int” arguments). The **every** loop (lines 2–18) reacts to each occurrence of **ENQUEUE**, creating and enqueuing the requested command, as illustrated in Figure 10:

- The initial state (box 0) assumes a pre-existing neutral ITEM in the root of the **waiting** field (④).
- Lines 4–11 assign a new subtree to the **tmp** field (box 1) with a new neutral ITEM (⑤, lines 5–6) linked to the set of commands to **MOVE** the turtle (⑥, lines 7–11).
- Lines 15–16 move the already **waiting** commands (③) to tail of the **tmp** node. The old location is automatically set to NIL (④). Note that we skip the neutral ITEM nodes of both **waiting** and **tmp** fields. This prevents the **waiting** root to become NIL and awake the ROOT node (*CODE-4*: 8) before we finish the enqueuing operation.
- Line 17 moves the **tmp** subtree (⑤) back to the **waiting** field, releasing the abandoned neutral ITEM, and notifying the ROOT node that the queue is no longer empty. The **tmp** field is automatically set to NIL. Note that the new **waiting** subtree preserves a neutral ITEM for subsequent enqueue operations.

## 4. RELATED WORK

Traversing data reactively in an imperative language requires dealing with concurrent updates. While attempts are made to make to this process more transparent, performance concerns ultimately require the programmer to specify behavior explicitly. In [4], one-way dataflow constraints are




Figure 10: Enqueuing new commands.

used to track updates in data structures in a reactive imperative language. Programmers need to annotate classes with constraint handler functions, which are then recursively called in the event of value updates. Another approach, focused on incremental computation, led to the development of self-adjusting computation [1], using a combination of dynamic dependency graphs and memoization. Keeping track of dependencies incurs significant overhead: traceable data types [2] are one approach for mitigating this issue by letting the programmer control the granularity of tracked data. In CÉU, tracking pointer updates is explicit (and enforced), and the move semantics for assignment of recursive data fields is designed to interact correctly with the traversal construct.

Introducing a control structure specifically for traversing recursive data is unusual, but parallels can be made to other languages. Generators trace their history back to CLU [11], but their implementation was stackless and therefore did not support recursion. Icon [5] allows yielding through recursive functions, but delegation is explicit via the `suspend` keyword. In [5], Icon generators are used for implementing goal-directed programming via “control backtracking” similarly to selector nodes presented in Section 3.1. Python originally introduced CLU-like stackless generators in version 2.2 [18], but those were later internally promoted to coroutines in order to support recursion [16]. Finally, in version 3.3, a form of delegation was introduced. In contrast, Lua [8] has first-class stackful coroutines: generator functions have to be constructed by wrapping coroutine objects, but there is no need for explicit delegation in recursive calls. CÉU builds its higher-level `traverse` construct on organisms, which are also a primitive for cooperative multitasking, supporting recursion and concurrency transparently.

## 5. CONCLUSION

We presented a new construct for traversing recursive data types incrementally, in the context of CÉU, an imperative reactive language with synchronous concurrency. The `traverse` construct encapsulates an idiom for performing recursive traversal by handling each step as a separate trail of execution. It allows parallel traversal using the language’s concurrency features and maintains its safety properties.

This kind of traversal can be done in CÉU through the use of organisms (pooled objects which launch their own execution trails) and orthogonal abortion via the `watching` construct. Combining these features to traverse a recursive data structure correctly, however, is not straightforward. Recursing in a way such that parallel constructs can be composed requires each step of the recursion to be a new execution trail. Ensuring that the traversal will not execute on a stale subtree in case the structure is modified requires the nodes

to be watched in order to perform abortions. Additionally, by presenting a control construct that is tied to a data structure, we can ensure bounded execution time, in line with the CÉU philosophy. By dealing with these concerns internally in the `traverse` statement, we make reactive traversal as easy to perform correctly as a recursive function call.

In the current implementation of recursive data types in CÉU, we impose restrictions to the kinds of structures that can be represented. The requirement of a tree hierarchy of ownership and move semantics for assignment of structure fields requires care in the design of algorithms manipulating these structures, as illustrated in Section 3.3. This is done to support static memory management with bounded memory pools for allocation and deterministic deallocation. Still, we do not feel that the restrictions are prohibitively limiting. For instance, persistent data structures in functional languages [12] operate under tighter design constraints.

Possibilities for future work are the introduction of type arguments for `data`, and investigating possibilities for relaxing CÉU pointer and reference semantics while maintaining their safety properties on recursive data.

## 6. REFERENCES

- [1] U. A. Acar et al. An experimental analysis of self-adjusting computation. *ACM Trans. Program. Lang. Syst.*, 32(1):3:1–3:53, Nov. 2009.
- [2] U. A. Acar et al. Traceable data types for self-adjusting computation. In *Proceedings of PLDI ’10, PLDI ’10*, pages 483–496. ACM, 2010.
- [3] O.-J. Dahl and K. Nygaard. SIMULA: an ALGOL-based simulation language. *Communications of the ACM*, 9(9):671–678, 1966.
- [4] C. Demetrescu et al. Reactive imperative programming with dataflow constraints. *ACM Trans. Program. Lang. Syst.*, 37(1):3:1–3:53, Nov. 2014.
- [5] R. E. Griswold, D. R. Hanson, and J. T. Korb. Generators in Icon. *ACM Trans. Program. Lang. Syst.*, 3(2):144–161, Apr. 1981.
- [6] C. Hecker. My liner notes for Spore. [http://chrishecker.com/My\\_liner\\_notes\\_for\\_spore](http://chrishecker.com/My_liner_notes_for_spore), 2009. Accessed: 2015-07-30.
- [7] J. Hoffmann and R. Brafman. Contingent planning via heuristic forward search with implicit belief states. In *Proc. ICAPS*, volume 2005, 2005.
- [8] R. Ierusalimsky. *Programming in Lua, Third Edition*. Lua.Org, 3rd edition, 2013.
- [9] D. Isla. Handling complexity in the Halo 2 AI. In *Game Developers Conference*, 2005.
- [10] M. N. Krohn et al. Events can make sense. In *USENIX Annual Technical Conference*, pages 87–100, 2007.

- [11] B. Liskov. A history of CLU. *SIGPLAN Not.*, 28(3):133–147, Mar. 1993.
- [12] C. Okasaki. *Purely Functional Data Structures*. Cambridge University Press, 1998.
- [13] S. Papert. *Mindstorms: Children, Computers, and Powerful Ideas*. Basic Books, Inc., 1980.
- [14] F. Sant’Anna et al. Safe system-level concurrency on resource-constrained nodes. In *SenSys’13*. ACM, 2013.
- [15] F. Sant’Anna et al. Structured synchronous reactive programming with Céu. In *Modularity’15*, 2015.
- [16] N. Schemenauer et al. PEP 255.  
<https://www.python.org/dev/peps/pep-0255/>.  
Accessed: 2015-08-03.
- [17] J. Slaney and S. Thiébaux. Blocks world revisited. *Artif. Intell.*, 125(1-2):119–153, Jan. 2001.
- [18] G. van Rossum and P. J. Eby. PEP 342.  
<https://www.python.org/dev/peps/pep-0342/>.  
Accessed: 2015-08-03.